

Week7 Learning Reflection

Name: Xinyu Chang

Time: 2022.05.09-05.15

Class: CSE 416 Intro. Machine Learning

Lectures: 13 and 14

Summary

1. In practice, SVD might be a useful way to compute eigenvectors during the PCA steps.
2. Precision and Recall tend to show a negative relationship.
3. Generally, for a given precision, we want recall to be as large as possible(vice versa)
 - a. one metric: area under the curve (AUC)
 - b. another metric: Set desired recall and maximize precision (precision at k)

Concepts

(Lecture 13)

Pre-class material

- Recommender Systems
 - recommends items to a user based on what we think will be the most “useful” for the user.
- Recommend System Challenges
 - explicit - user tells us what she likes
 - implicit - we try to infer what she likes from usage data
 - cold start problem.
 - solution: side information
 - top-k vs. diverse recommendations
 - top-k: redundant
 - diverse: users are multi-faceted and we want to hedge our bets
- Popularity Model
 - Simplest Approach: Recommend whatever is popular
 - Limitation: No personalization; Feedback loops
- Classification Model
 - Train a classifier to learn whether or not someone will like an item
 - Pros:
 - personalized;
 - features can capture contexts;
 - can even handle limited user history
 - Cons:
 - features might not be available or hard to work with;
 - Often doesn't perform well in practice when compared to more advanced techniques like collaborative filtering.
 - can still lead to feedback loops
- Feedback loops

In-class Material

- Dimensionality reduction: the task of representing the data with a fewer number of dimensions, while keeping meaningful relations between data
 - example: embedding pictures.
 - embed high dimensional data in low dimensions to visualize the data.
 - Goal: Similar images should be near each other
 - Easier Learning: fewer parameters, no curse of dimensionality
 - Hard to visualize more than 3D
 - Discover intrinsic dimensionality of the data
 - high dimensional data can sometimes be truly lower dimensional.
- Projection
 - Linear projection
 - Project data into 1 dimension along a line
 - Could do something like feature importance and find the subset of features with the most meaningful information
 - A popular approach is to create new features that are combinations of existing features.
 - Benefit: We can do this in the unsupervised setting when we only know x and not y.
- Principal Component Analysis (PCA)
 - Idea: Use a linear projection from d-dimensional data to k-dimensional data
 - Algorithm
 - Input data: An n*d data matrix X: each row is an example
 - 1. Recenter Data: Subtract mean from each row

$$X_c \leftarrow X - \bar{X}[1:d]$$

- 2. Compute spread/orientation: Compute covariance matrix Σ

$$\Sigma[t, s] = \frac{1}{n} \sum_{i=1}^n x_{c,i}[t] x_{c,i}[s]$$

- 3. Find basis for orientation: Compute eigenvectors of Σ
 - Select k eigenvectors u_1, \dots, u_k with largest eigenvalues
- 4. Project Data: Project data onto principal

$$z_i[1] = u_1^T x_{c,i}$$

...

$$z_i[k] = u_k^T x_{c,i}$$

- Choose the projection that minimizes reconstruction error
 - Idea: The information lost if you were to undo the projection
- Reconstruction error
 - Reconstruction: Reconstruct original data only knowing the projection

$$\hat{x}_i[1:d] = \bar{X}[1:d] + \sum_{j=1}^k z_i[j] u_j$$
 - error: $\|\hat{x}_i[1:d] - x_i\|_2^2$
- How to choose the line direction that minimizes reconstruction error
- Geometric understanding of PCA
- Eigenvectors / Eigenfaces
 - Note: We don't particularly care about the math/algorithms to find eigenvectors. Covered in linear algebra classes :)
- When PCA doesn't work
 - PCA assumes there is a low dimensional linear subspace that represents the data set.
 - May want to look into non-linear dimensionality reduction
 - Manifold learning
 - Popular: SDD Maps, Isomap, LLE, t-SNE
 - Covariance matrix can be very large with high-dimensional data
 - this means finding the principal components will be slow
 - In practice, can use Singular Value Decomposition
 - can be used to find the k eigenvectors with largest eigenvalue
 - very fast implementations

Section material

- Matrix Completion
 - is an impossible task without some assumptions on data.
 - $Rating(u, v) = \hat{L}_u \cdot \hat{R}_v$
 - Quality metric
 - $\hat{L}, \hat{R} = \underset{L, R}{argmin} \sum_{u, v: \text{inequ}} (L_u * R_v - r_{uv})^2$
- Coordinate Descent
 - Issue: No unique solution for determining L and R
 - Use coordinate descent to optimize for one coordinate at a time
 - If the system underdetermined, use regularization
- Cold Start and Blended Model
 - Issue: we have a new user and want to give personalized recommendations based on their preferences

- Solution: Blended Model/Featurized Matrix Factorization

(Lecture 14)

Pre-class material

- Co-occurrence matrix
 - Idea: People who bought this, also bought...
 - $C \in R^{m \times m}$ (m is the number of items)
 - C_{ij} = # of users who bought both item i and j
 - Recommend people things with largest counts
 - Issues:
 - Popular items might drown out other useful items.
 - Solutions
 - Normalization: The count matrix C needs to be normalized. Normalize the counts by using the Jaccard similarity.
 - Could also use something like Cosine similarity, but Jaccard is popular.
 - Issue:
 - What if I know the user has bought diapers and milk?
 - Solution:
 - Take the average similarity between items they have bought.
 - Or can put different weight,
 - then find the item with the highest average score.
 - Pros
 - It is personalized to user
 - Cons
 - Does not utilize context (time of day), user features (age), product features; scalability (m^2); cold start problem.
- Jaccard similarity- capture how often the co-occurrence appears
 - $S_{ij} = \frac{\# \text{purchased } i \text{ and } j}{\# \text{purchased } i \text{ or } j} = \frac{C_{ij}}{T_i + T_j - C_{ij}}$ where T_i is the number of times people buy i in total
- Matrix factorization
 - Need Assumptions: e.g. Assume there are k types of movies, which users have various interests in.
 - User factors
 - e.g. how much the user likes each of the categories.
 - Item factors
 - e.g. how much the movie belongs to each of the categories.

In-class Material

- Matrix factorization quality metric

- Find L and R that when multiplied, achieve predicted ratings that are close to the values that we have data for.
- $\hat{L}, \hat{R} = \underset{L, R}{\operatorname{argmin}} \sum_{u, v: r_{uv} \neq ?} (L_u \cdot R_v - r_{uv})^2$
- $L_u \cdot R_v$ - predicted rating; r_{uv} - actual rating
- Coordinate descent
 - i.e. Alternating Least Squares
 - implement CD to find \hat{L} and \hat{R}
 - Gradient Descent is not used much in practice to optimize this.
 - Goal: Minimize some function $g(w) = g(w_0, w_1, \dots, w_D)$
 - Instead of finding optima for all coordinates, do it for one coordinate at a time,
 - Guaranteed to find an optimal solution under some constraints
 - Algorithm:
 - Initialize $\hat{w} = 0$ (or smartly)
 - while not converged:
 - pick a coordinate j
 - $\hat{w}_j = \underset{w}{\operatorname{argmin}} g(\hat{w}_0, \dots, \hat{w}_{j-1}, w, \hat{w}_{j+1}, \dots, \hat{w}_D)$
 - \uparrow here, we keep all other w 's except w_j as constant.
- Coordinate descent for matrix factorization
 - Fix movie factors R and optimize for L_u
 - Independent least squares for each user (n optimizations)
 - $V_u = \text{all movies user } u \text{ has rated}$
 - $\underset{L_1, \dots, L_n}{\operatorname{argmin}} \sum_{u, v: r_{uv} \neq ?} (L_u \cdot R_v - r_{uv})^2$
 - $= \underset{L_1, \dots, L_n}{\operatorname{argmin}} \sum_u \sum_{v \in V_u} (L_u \cdot R_v - r_{uv})^2$
 - changing of the factors for one user, only will affect the predictions of that one user.
 - Solve for each user u :
 - $\hat{L}_u = \underset{L_u}{\operatorname{argmin}} \sum_{v \in V_u} (L_u \cdot R_v - r_{uv})^2$
 - Looks like linear regression:
 - $\underset{w}{\operatorname{argmin}} \sum_{i=1}^n (w^T h(x_i) - y_i)^2$
 - use Gradient Descent: with k coefficients for one user.
 - Fix user factors, and optimize for movie factors separately
 - Independent least squares for each movie (m optimizations)
 - $\hat{R}_v = \underset{R_v}{\operatorname{argmin}} \sum_{u \in U_v} (L_u \cdot R_v - r_{uv})^2$
 - Repeatedly do these steps until converge to local optima

- Underdetermined -> use regularization.
- $nk + mk$ parameters we will have to learn.
- Topics
 - The features found by matrix factorization don't always correspond to something meaningful
 - The exact values are meaningless but the directions mean something
- Blending models: Learn a model to supplement the matrix factorization model
 - Cold Start problem exists.
 - create a feature vector for each movie
 - Define weights on these features for all users, global.
 - $w_G \in R^d$
 - Fit linear model
 - $\hat{r}_{uv} = w_G^T h(v)$
 - $\hat{w}_G = \underset{w}{\operatorname{argmin}} \sum_{r_{u,v}: r_{uv} \neq ?} (w^T h(v) - r_{uv})^2 + \lambda ||w||$
 - Prob: feedback loop exists. No general solution, needs to depend on context. Relatively new area need to explore.
 - Featurized matrix factorization
 - Feature-based approach
 - Feature representation of user and movie fixed
 - Can address cold start problem
 - Matrix factorization approach
 - Suffers from cold start problem
 - User and Movie features are learned from data
 - A unified model
 - $\hat{r}_{uv} = f(\hat{L}_u \cdot \hat{R}_v, (w_G + w_u)^T h(u, v))$
- Evaluating recommender systems - precision/recall
 - We want to look at our set of recommendations and ask:
 - How many of our recommendations did the user like? Precision
 - How many of the items that the user liked did we recommend? Recall
 - Precision
 - $\frac{\# \text{ liked and shown}}{\# \text{ shown}}$
 - Recall
 - $\frac{\# \text{ liked and shown}}{\# \text{ liked}}$

Uncertainties

1. Why do we want to find the PCA direction that will maximize the variance along with it?
 - a. More variance, so that it will show the differences clearer.